

# Class08

Samuel Fisher (A18131929)

## Save your input data file into your Project directory

```
fna.data <- "WisconsinCancer.csv"
```

## Complete the following code to input the data and store as wisc.df

```
wisc.df <- ____(fna.data, row.names=1)
```

## Exploratory Data Analysis

First we load the Wisconsin cancer dataset and prepare it for analysis.

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)  
head(wisc.df, 4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
					8.589
					3.398

84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
	fractal_dimension_worst				
842302		0.11890			
842517		0.08902			
84300903		0.08758			
84348301		0.17300			

We remove the diagnosis column so unsupervised methods do not use the known labels.

```
wisc.data <- wisc.df[, -1]
```

## Diagnosis Vector

We save the diagnosis column as a factor for later comparison and plotting.

```
diagnosis <- as.factor(wisc.df$diagnosis)
```

## Q1 - Number of observations

We check how many observations (rows) are in the dataset.

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in the dataset

## Q2 - Number of malignant samples

We want to count how many samples are labeled malignant in the diagnosis vector.

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

212 observations are malignant

## Q3 - Number of \_mean features

We want to count how many variable names end with \_mean.

```
length(grep("_mean$", colnames(wisc.data)))
```

```
[1] 10
```

10 variables names end with \_mean

## Principal Component Analysis

### Check column means and standard deviations

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst

6.573234e-02

6.186747e-02

1.806127e-02

The variables have very different standard deviations, so scaling is required before performing PCA.

## PCA Model

**Perform PCA on wisc.data by completing the following code**

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
```

We want to look over the PCA summary to see how much variance of each principal component

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

## **Q4 - Variance found by PC1**

Proportion of Variance — PC1 = 0.4427, therefore PC1 captures 44.27% of the total variance.

## Q5 - PCs needed for 70% variance

PC1 = 0.4427 PC2 = 0.6324 PC3 = 0.72636 ← first value greater than or equal to 0.70

Three principal components are needed to explain at least 70% of the variance.

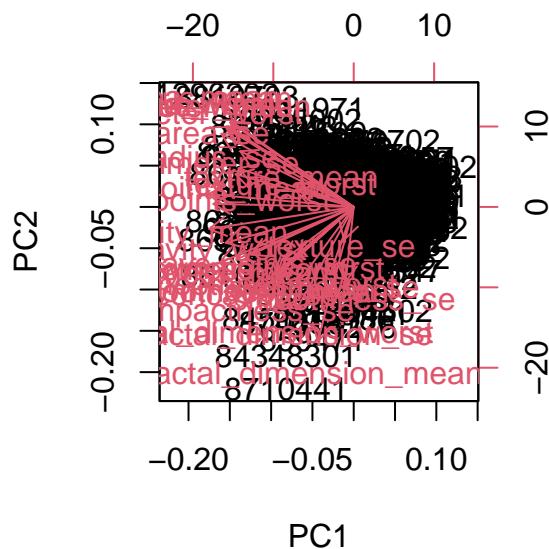
## Q6 - PCs needed for 90% variance

$\text{PC6} = 0.88759$   $\text{PC7} = 0.91010 \leftarrow$  first value greater than or equal to 0.90 Seven principal components are needed to explain at least 90% of the variance.

# Interpreting PCA Results

We want to create a PCA biplot to visualize both sample scores and feature loadings.

```
biplot(wisc.pr)
```



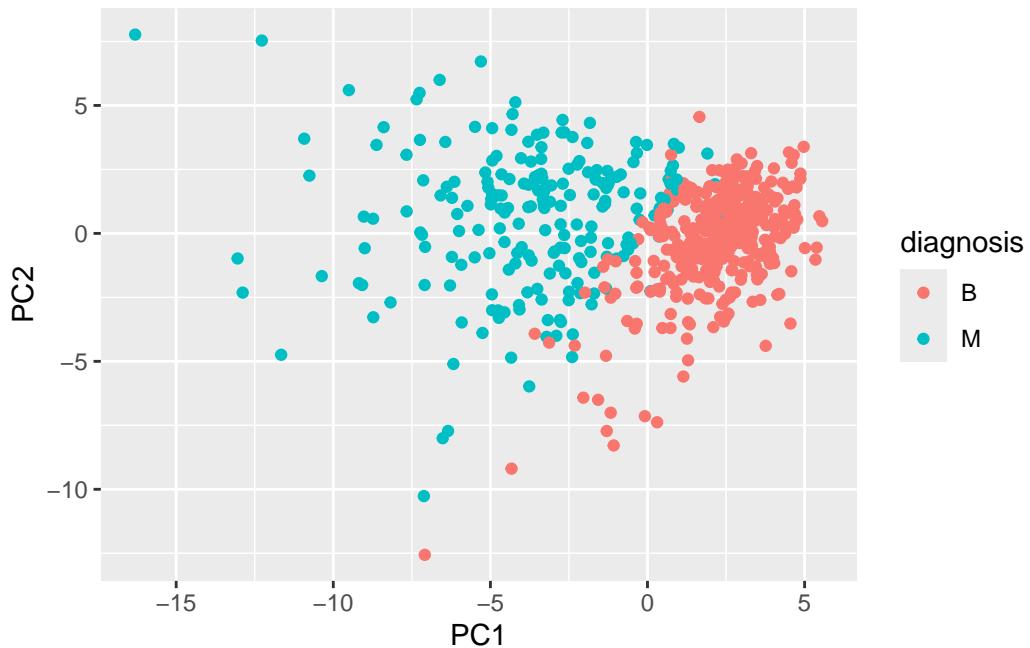
## Q7 - Biplot Interpretation

The biplot is incredibly difficult to interpret because it is chaotic and with many things overlapping. The row labels are all over the figure, making patterns and group separation hard to see and understand.

### PC1 vs PC2 Plot

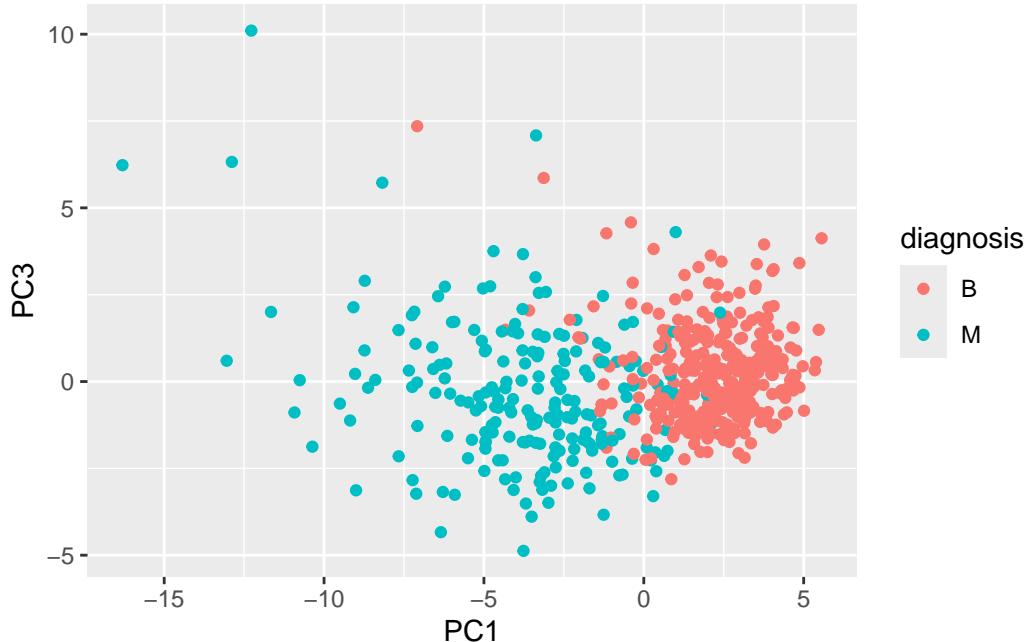
Scatter plot observations by components 1 and 2

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col = diagnosis) +
  geom_point()
```



### Q8 - Compare plots of PC1 vs PC2 to plot of PC1 vs PC3

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col = diagnosis) +
  geom_point()
```



## PC1 vs PC3 Plot

The PC1 vs PC2 plot has a clearer separation between M and B samples than the PC1 vs PC3 plot. This shows that PC2 captures more class-separating structure than PC3. PC1 seems to drive most of the separation overall, while later components add less discriminatory power.

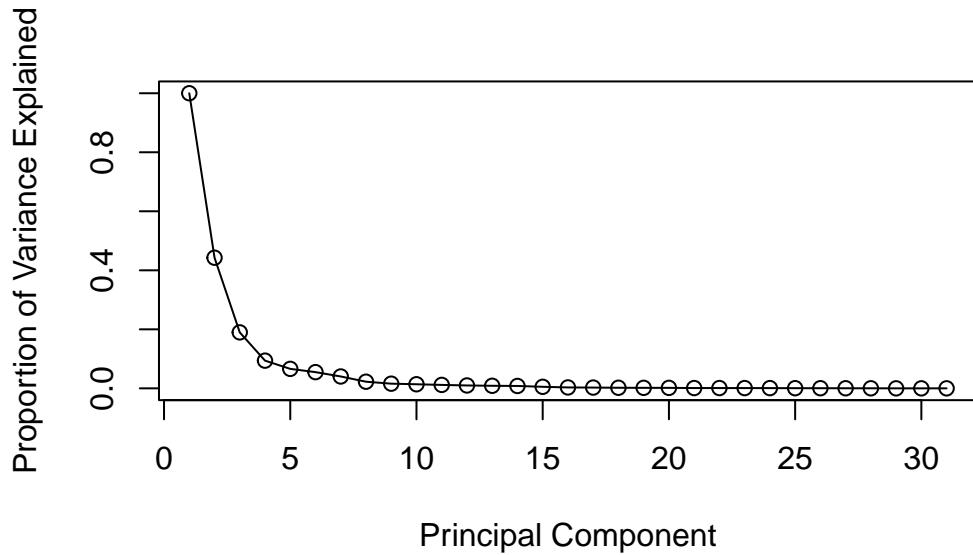
### Variance Explained

### Calculate variance of each component

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

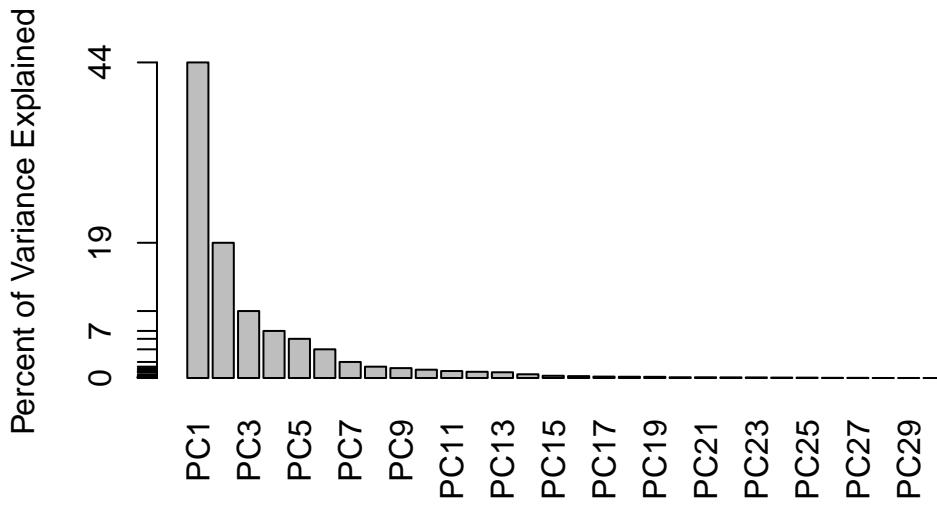
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var / sum(pr.var)
plot(c(1,pve), xlab = "Principal Component",
ylab = "Proportion of Variance Explained",
ylim = c(0, 1), type = "o")
```



**Alternative scree plot of the same data, note data driven y-axis**

```
barplot(pve, ylab = "Percent of Variance Explained",
names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



## Q9 — Plot Interpretation

```
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

```
sort(wisc.pr$rotation[,1], decreasing=TRUE) [1:5]
```

smoothness_se	texture_se	symmetry_se
-0.01453145	-0.01742803	-0.04249842
fractal_dimension_mean	fractal_dimension_se	
-0.06436335	-0.10256832	

```
sort(abs(wisc.pr$rotation[,1])), decreasing=TRUE) [1:5]
```

concave.points_mean	concavity_mean	concave.points_worst
0.2608538	0.2584005	0.2508860
compactness_mean	perimeter_worst	
0.2392854	0.2366397	

The loading value for concave.points\_mean in PC1 is 0.2608538. There are no features with a larger absolute loading than this one. It is the largest contributor to PC1 (with concavity\_mean and concave.points\_worst being slightly smaller but similar).

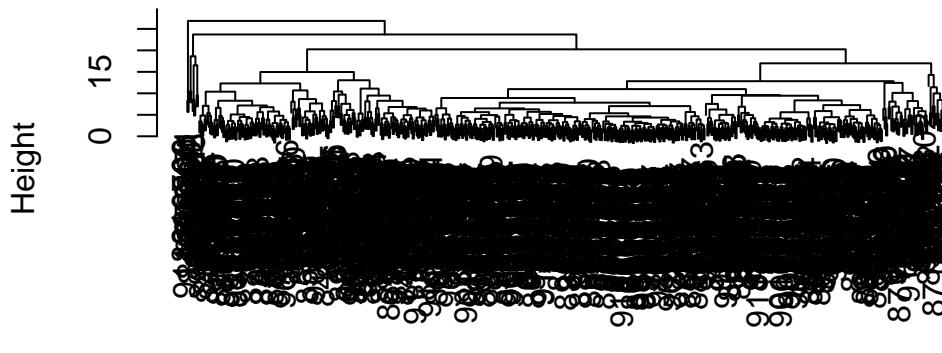
## Hierarchical Clustering

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
```

Perform hierarchical clustering using complete linkage and plot

```
wisc.hclust <- hclust(data.dist, method = "complete")
plot(wisc.hclust)
```

## Cluster Dendrogram

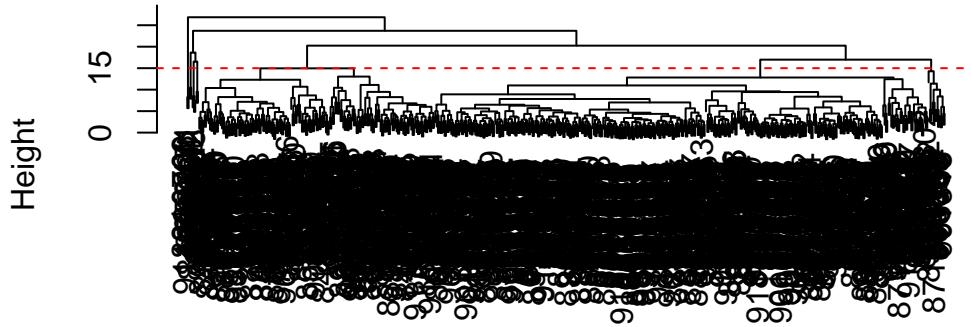


data.dist  
hclust (\*, "complete")

## Q10

```
plot(wisc.hclust)
abline(h = 15, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

The height at which the clustering model has 4 clusters is approximately 15.

```
wisc.clusters <- cutree(wisc.hclust, k = 4)
```

```
table(wisc.clusters, diagnosis)
```

wisc.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

## Q12

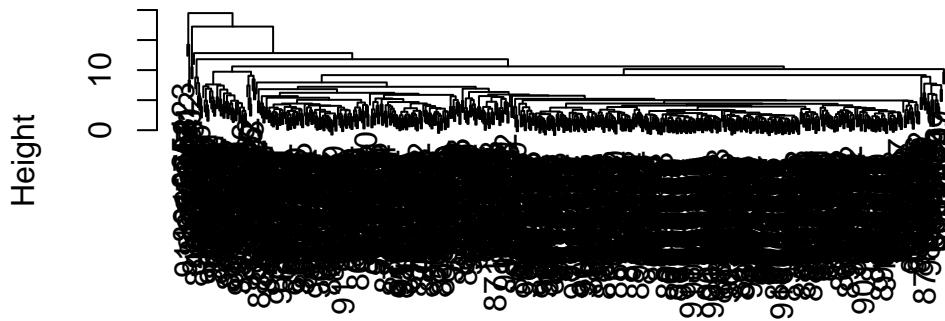
Ward.D2 gives the best results because it creates the cleanest and most balanced cluster separation while minimizing cluster variance within, which fits this dataset well.

## Clustering on PCA Results

We will do the hierarchical clustering again but using average linkage to compare results.

```
wisc.hclust.avg <- hclust(data.dist, method = "average")
plot(wisc.hclust.avg)
```

## Cluster Dendrogram



data.dist  
hclust (\*, "average")

```
wisc.clusters.avg <- cutree(wisc.hclust.avg, k = 2)
table(wisc.clusters.avg, diagnosis)
```

	diagnosis	
wisc.clusters.avg	B	M
1	357	209
2	0	3

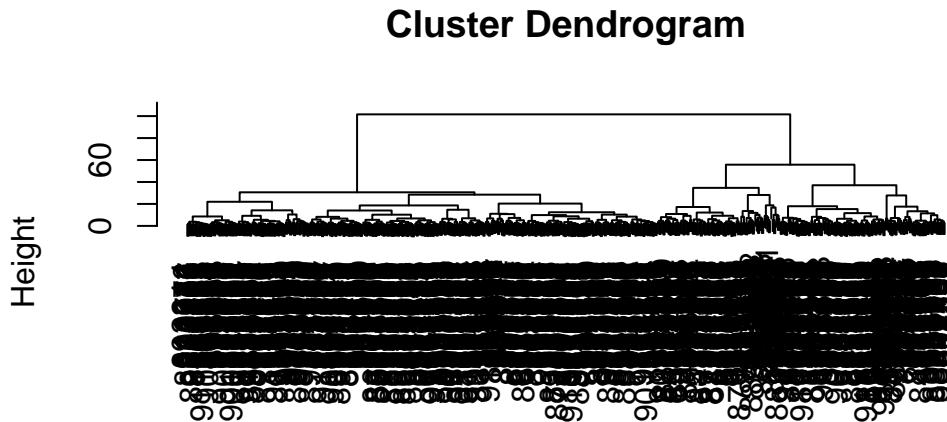
## Average Linkage Cluster Comparison

Using average linkage makes a similar result to complete linkage. The clusters still don't align perfectly with diagnosis labels. This shows that hierarchical clustering w/o labels doesn't perfectly separate B and M samples.

## Clustering on PCA Scores

### Hierarchical clustering on PCA scores (first 7 PCs)

```
pc.dist <- dist(wisc.pr$x[,1:7])
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method = "ward.D2")
plot(wisc.pr.hclust)
```



```
dist(wisc.pr$x[, 1:7])
hclust (*, "ward.D2")
```

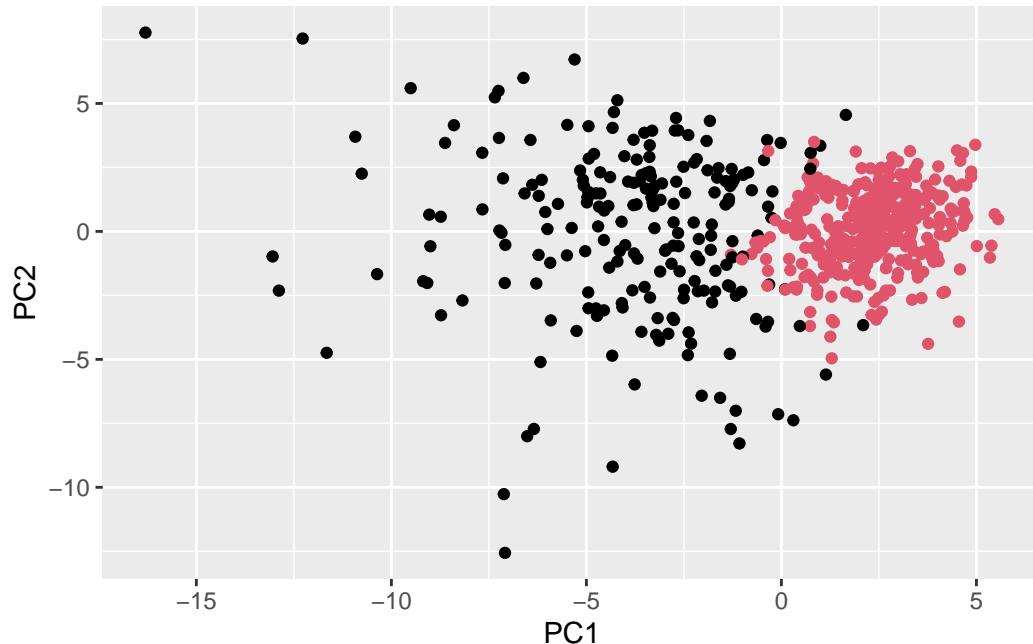
```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
216 353
```

```
table(grps, diagnosis)
```

grps	B	M
1	28	188
2	329	24

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC2) +  
  geom_point(col=grps)
```



Use the distance along the first 7 PCs for clustering i.e. wisc.pr\$x[, 1:7]

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method="ward.D2")
```

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Q13

Compare to actual diagnoses

```
table(wisc.pr.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.pr.hclust.clusters   B   M
                           1 28 188
                           2 329 24

```

It separates pretty well, with 52 total misclassified. 28 Benign mixed into the mostly malignant cluster plus 24 malignant mixed into the mostly-benign cluster.

## Q14

```
table(wisc.clusters, diagnosis)
```

```

diagnosis
wisc.clusters   B   M
               1 12 165
               2  2  5
               3 343 40
               4  0  2

```

Clustering on the PCA transformed data separates the diagnoses better than clustering on the original features. The PCA based model generates two better clusters that have less mixed benign/malignant cases. The original data clustering spreads samples across more mixed clusters.

## Prediction

```

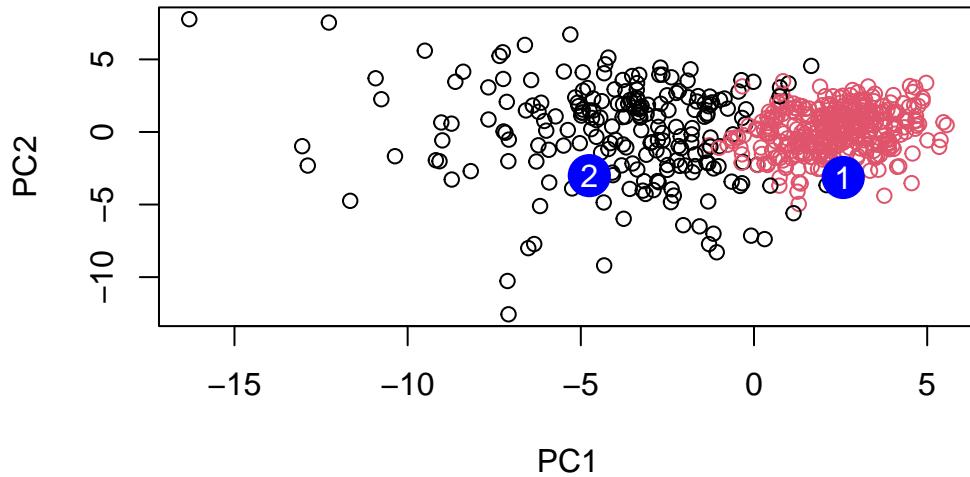
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc

```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882

```
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
      PC15      PC16      PC17      PC18      PC19      PC20
[1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
      PC21      PC22      PC23      PC24      PC25      PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27      PC28      PC29      PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



## Q16

Patient 1 should be prioritized for a follow up because it falls within the malignant like cluster. Patient 2 groups with the benign like cluster.